

Problem A. Array

Let $prev_i$ be the position of previous occurrence of a_i (set as -1 when there is no such position) and $next_i$ be the position of next occurrence of a_i (set as $n + 1$ when there is no such position). Apparently, there's no need to change such a_x that $prev_x \neq -1$. So, we can only focus on those a_x where $prev_x = -1$. In addition, it is easy to observe changing a_x into some a_y with $prev_y = -1$ is not worse than the other choices.

1 change a_x into some a_y ($y < x$)

The cost will be $|a_x - a_y| + \frac{x(x-1)}{2} - \frac{next_x(next_x-1)}{2}$. We should make $|a_x - a_y|$ as minimal as possible. It can be done by using a set to maintain all the value before x . We can use a disjoint union set or even a bidirectional list as well.

2 change a_x into some a_y ($x < y$)

The cost will be $|a_x - a_y| + \frac{y(y-1)}{2} - \frac{next_x(next_x-1)}{2}$. Assume $a_x \leq a_y$ ($a_x \geq a_y$ is the same), this can be done by using a segment tree to maintain the minimum value of $\frac{y(y-1)}{2} - a_y$. Since the type of query is fairly simple, we can use a fenwick tree too.

Problem B. Chiaki Chain

Some observations of the k -th order Chiaki Chain with n vertices and m edges:

- $n + k - 1 = m$;
- connected, no loops, no multiple edges;
- each vertex in a cycle has degree 2 except one vertex connected with the sub-chain;
- after removing all the cycles and sub-chains, the remainder is a chain.

Firstly, use a simple depth-first search to find all the cycles and check whether the degree constraint is satisfied.

Then remove the cycles and sub-chains greedily using breadth-first search and we should get a chain or a single vertex (all other graphs we get is invalid).

If we get a single vertex, we should check if the length of sub-chain is greater than 1 when $k = 1$ and check if the length of at least one sub-chain is greater than 1 when $k = 2$ and check if the length of at least two sub-chains is greater than 1 when $k = 3$. In this case, $k > 3$ is invalid.

Otherwise, we should enumerate each vertex with degree 3 and check if the sub-chain is valid: the length of the sub-chain is at least 1. In addition, if the vertex is the end of the chain and two sub-chains extend from it, the length of at least one sub-chain should be greater than 1.

Problem C. Cut the Plane

At first, find a line to split the points into two parts — the first part has $\lfloor \frac{n}{2} \rfloor$ points and the second has $\lceil \frac{n}{2} \rceil$.

Then, find a tangent line of the convex hulls of two parts and we can use a line to separate both the tangent points from these two parts. Repeat until there is only one point left in each part.

Problem D. Edges Counting

Let t_n be the number of trees with n labeled vertices, and its exponential generating function be $T(x) = \sum_{n \geq 1} \frac{t_n}{n!} x^n$. Cayley's formula states that $t_n = n^{n-2}$. Similarly, let rt_n be the number of rooted trees with n labeled vertices, and its exponential generating function be $RT(x) = \sum_{n \geq 1} \frac{rt_n}{n!} x^n$. Since $rt_n = n \cdot t_n$, we can figure out that $RT(x) = xT'(x)$, where $T'(x)$ is the derivative function of $T(x)$ with respect to x .

We call a simple graph generated from a tree with an additional edge "pseudotree", which contains exactly one simple cycle, and if one removes the cycle from the pseudotree, the graph will be changed into a forest. Let p_n be the number of pseudotrees with n labeled vertices, and its exponential generating function be $P(x) = \sum_{n \geq 1} \frac{p_n}{n!} x^n$. An observation is that $P(x) = \sum_{k \geq 3} \frac{RT^k(x)}{k!} \frac{(k-1)!}{2} = \sum_{k \geq 3} \frac{RT^k(x)}{2k}$. The combinatorial meaning is to enumerate the number of vertices on the only cycle, denoted by k , and then the remain parts are k rooted trees, so we can get the conclusion by convolution. Moreover, we can get $P'(x) = \frac{RT'(x)}{1-RT(x)} \frac{RT^2(x)}{2}$, which may help calculate p_1, p_2, \dots, p_n .

After that, we can count the number of good graphs. That is to say, let g_n be the number of good graphs with n labeled vertices, and its exponential generating function be $G(x) = \sum_{n \geq 1} \frac{g_n}{n!} x^n$, and we can enumerate the number of components, denoted by k , set k unordered components and get $G(x) = \sum_{k \geq 0} \frac{(T(x)+P(x))^k}{k!} = e^{T(x)+P(x)}$. By the way, it can be rewritten as $G'(x) = G(x)(T(x) + P(x))'$. (Is that easier for computing?)

Now let's count the number of edges on cycles. If we know the number of good graphs, we could count the contribution of each component. More precisely speaking, we may enumerate a component, which should be a pseudotree, and count the number of vertices on its cycle. The generating function should be modified as $\sum_{k \geq 3} k \frac{RT^k(x)}{2k} = \frac{RT(x)}{1-RT(x)} \frac{RT^2(x)}{2}$. Let the generating function be $C(x)$, and we know $G(x) \cdot C(x)$ yields the answer.

Due to that the modulus may not be a prime larger than n , so we can hardly calculate each generating function. However, each integer sequence (e.g. $t_n, rt_n, p_n, g_n, \dots$) is easy to get. One can use the binomial coefficients to calculate the convolution. The time complexity is $\mathcal{O}(\max^2\{n\})$.

Problem E. Equanimous

It should be the hardest problem in this contest if you haven't seen its weakened version. The conclusion is that we can construct a deterministic finite automaton (DFA), whose minimal automaton contains only 715 states, and use digit DP on this DFA. To improve the query time complexity, we can prepare more information at first and calculate less information for each query. The time complexity could be $\mathcal{O}(D^2LS + TL)$, where $D = 10$ is the radix, $L = 100$ is the maximal length of a number, $S = 715$ is the number of states on DFA, and T is the number of queries.

To build the DFA, we need to understand how to calculate $f(m)$ for a given integer m firstly. Let $dp(i, j)$ be the possibility (true or false) such that the first i digits of m could form an expression whose absolute value is exactly j . The transition should be $dp(i-1, j) \rightarrow dp(i, j+d_i), dp(i-1, |j-d_i|)$, where d_i is the i -th digits of m . Actually, we don't need to memorize the information of $j \geq K$ for some integer K . We know $0 \leq f(m) \leq 9$, and the worst case for this DP occurs when m contains consecutive digits like $\underbrace{99 \dots 9}_{8 \text{ times}} \underbrace{88 \dots 8}_{9 \text{ times}}$, so we can just set K as $(8 \times 9 + 1)$.

The processing of such dynamic programming builds the connection among $d_1, 10d_1+d_2, 100d_1+10d_2+d_3, \dots, m$. After removing less useful information with $j \geq K$, the information of each integer (for the sake of transition) could be represented as a 01-sequence of length 73. Actually, the number of different sequences is fairly small, so we can build a DFA consisting of these sequences and corresponding transitions. Inspired by the aforementioned worst case, we can minimize the DFA using partition refinement through at most 9 iterations.

The remaining part is just a typical digit DP problem. We can calculate $DP(i, j, k)$ as the number of integers m of length i such that, if we start from the j -th state on the DFA and pass through the digits

of m , we can reach a state with $f(m') = k$. Besides, calculate $S(i, j, k, l)$ as $\sum_{x=0}^l DP(i-1, \text{go}(j, x), k)$, where $\text{go}(j, x)$ represents the next state from the j -th state passing through one more digit x . Calculating S may help calculating DP and handle queries fast.

Problem F. Fighting Against Monsters

This problem can be solved case by case. In the following discussion, the time complexity should be $\mathcal{O}(4H^3)$, where $H = \max\{HP_A, HP_B\}$.

1 $HP_C \leq 5050$

In this case, battles should end before or at 102-th second, no matter these monsters die in which order. Let $dp(i, j, k)$ be the minimum total damages the hero need to suffer at the end of the i -th second, when the first monster has suffered total damages of value j and the second monster has suffered total damages of value k (both including the damages overflow). The transition is easy since one can conclude the boss monster has suffered total damages of value $\left(\frac{i(i+1)}{2} - j - k\right)$.

2 $HP_C > 5050$

2.1 The boss monster dies at first

In this case, we observe that, after killing the boss monster, we can kill each alive monster in one second, no matter how many damages they have suffered before the boss monster dies. So, we can solve the case using a greedy algorithm.

2.2 The boss monster dies in the middle

In this case, we can ignore the last died monster first. Let $T(x)$ be the lowest integer t such that $\frac{t(t+1)}{2} \geq x$. We can prove that the first died monster, denoted by p , could die at the end of $T(HP_p)$ -th second, and thus the boss monster could die at the end of $T(HP_p + HP_C)$ -th second, which is an optimal strategy.

2.3 The boss monster dies at the end

In this case, the little monsters should die before or at 101-th second. One can just apply DP similar to the aforementioned one, enumerate the total damages of these two little monsters and then determine the time the boss monster dies.

Problem G. Mysterious Triple Sequence

It's a bit hard to solve the problem directly. Let's reduce this problem into the **discrete logarithm** problem which is easier to solve in **sqrt** time.

1 A series of lemmas

Lemma 1. Let $f_0 = 0, f_1 = 1, f_{n+2} = 2f_{n+1} + f_n$ ($n \in \mathbb{N}^+$). It shows that $(a_k, b_k, c_k) = (f_{2^{k+1}}, f_{2^k}, f_{2^{k-1}})$.

Proof. There are many approaches to prove that (and further things), so I'd like to leave a type of brief introduction here.

Let $F = \begin{bmatrix} 0 & 1 \\ 1 & 2 \end{bmatrix}$. It can be proved by induction that $F^n = \begin{bmatrix} f_{n-1} & f_n \\ f_n & f_{n+1} \end{bmatrix}$ ($n \in \mathbb{N}^+$).

There is an observation that $\begin{bmatrix} f_{n-1} & f_n \\ f_n & f_{n+1} \end{bmatrix} \begin{bmatrix} f_{m-1} & f_m \\ f_m & f_{m+1} \end{bmatrix} = \begin{bmatrix} f_{n-1}f_{m-1} + f_nf_m & f_{n-1}f_m + f_nf_{m+1} \\ f_nf_{m-1} + f_{n+1}f_m & f_nf_m + f_{n+1}f_{m+1} \end{bmatrix}$ which is helpful to prove $(a_k, b_k, c_k) = (f_{2^k+1}, f_{2^k}, f_{2^k-1})$ by induction. \square

Lemma 2. Sequence $\{(f_{k+1}, f_k)\}_{k=0}^{\infty}$ is a pure periodic sequence in modulo any integer p .

Proof. Assuming that there exist two distinct integers x and y such that $(f_{x+1}, f_x) \equiv (f_{y+1}, f_y) \pmod{p}$, it can be proved that $(f_{x+k+1}, f_{x+k}) \equiv (f_{y+k+1}, f_{y+k}) \pmod{p}$ for any positive integer k just by using matrix multiplication of matrix F .

Furthermore, it is easy to show that $F^{-1} = \begin{bmatrix} -2 & 1 \\ 1 & 0 \end{bmatrix}$ always exists in modulo any integer p and thus it is possible to extend the above k from any positive integer ($k \in \mathbb{N}^+$) to any integer ($k \in \mathbb{Z}$). \square

Lemma 3. Let $L(p)$ be the minimal positive period of sequence $\{(f_{k+1}, f_k)\}_{k=0}^{\infty}$ in modulo integer p . It shows that $L(p) \leq \frac{8}{3}p$ for any integer p .

Proof. If we replace sequence $\{f_n\}$ by Fibonacci sequence, it will be a classic conclusion that $L(p) \leq 6p$. In addition, it is not difficult to conclude the formula of the period for $\{f_n\}$ by a similar approach. If you want to know more details, please check “Pisano period - Wikipedia” and “The Period of the Fibonacci Sequence Modulo j ”. If not, the following are conclusions.

Case 1 If $p = 2$, $L(p) = 2$.

Case 2 If p is an odd prime such that there exists an integer x satisfying $x^2 \equiv 2 \pmod{p}$, $L(p)$ is a divisor of $(p-1)$. In that case, $p \equiv 1$ or $7 \pmod{8}$.

Case 3 If p is an odd prime such that there exists no integer x satisfying $x^2 \equiv 2 \pmod{p}$, $L(p)$ is a divisor of $2(p+1)$. In that case, $p \equiv 3$ or $5 \pmod{8}$.

Case 4 If $p = q^e$ such that q is a prime and e is an integer greater than 1, $L(p)$ is a divisor of $q^{e-1}L(q)$.

Case 5 If $p = q_1^{e_1}q_2^{e_2}\cdots q_m^{e_m}$ such that q_1, q_2, \dots, q_m are distinct primes, $L(p)$ equals to the least common multiple of $L(q_1^{e_1}), L(q_2^{e_2}), \dots, L(q_m^{e_m})$.

These cases are easy to prove using Fermat’s little theorem and Euler’s criterion, so we leave them to the reader.

The worst case such that $\frac{L(p)}{p}$ is maximum occurs when $p = q^e$, q is a prime, e is an integer and $q \equiv 3$ or $5 \pmod{8}$. In that case, $\frac{L(p)}{p} \mid \frac{2(q+1)q^{e-1}}{q^e} = \frac{2q+2}{q} \leq \frac{8}{3}$. \square

By the way, please note that $L(181^4) = 2158425724$, which is actually the only one case such that $p \leq 2^{30}$ but $L(p) \geq 2^{31}$ (random tests might not cover this case), so be careful with the 32-bit integer.

2 Solution

We can reduce the problem into two stage:

1. First, find the integer u ($0 \leq u < L(p)$) such that $(x, y, z) \equiv (f_{u+1}, f_u, f_{u-1}) \pmod{p}$. Because $\{(f_{k+1}, f_k)\}_{k=0}^{\infty} \pmod{p}$ is pure periodic, you can use “Baby-step giant-step algorithm” (a type of meet-in-the-middle algorithm) directly.
2. Then, find the integer k ($k \geq m$) such that $2^k \equiv u \pmod{L(p)}$. Because $\{2^k\}_{k=0}^{\infty}$ is not pure periodic in modulo $L(p)$ ($L(p)$ is always even), you should enumerate the acyclic part and reduce $L(p)$ into an odd number and then use Baby-step giant-step algorithm.

However, solution in time complexity $\mathcal{O}(n\sqrt{p}\log p)$ or $\mathcal{O}(n\sqrt{p})$ would be rejected. Actually, this type of meet-in-the-middle algorithm could be optimized easily for multi-queries. Let's set a threshold T first (we will determine an optimal value for it soon).

In the first stage, we need to find the least power of the invertible matrix F which equals to a given matrix G . Assuming that $F^{xT+y} = G$ ($0 \leq y < T$), we have $F^y = G(F^{-T})^x$. After precalculating F^0, F^1, \dots, F^{T-1} , you can just enumerate x and check if y exists (by binary search or hash). The time complexity for multi-queries is $\mathcal{O}(T \log T + n \frac{L(p)}{T} \log T)$. (note: there are only two element in F^y is necessary to memorize)

In the second stage, we need to find the least power of 2 which equals to a given number u . let $L(p)$ be $2^\alpha L'(p)$ such that $L'(p)$ is an odd integer. If there exists integer k such that $0 \leq k < \alpha$, $2^k \equiv u \pmod{L(p)}$, that is finished. If not, we can reduce the problem into $2^{k-\alpha} \equiv \frac{u}{2^\alpha} \pmod{L'(p)}$ (the minimal positive period of 2 will be a divisor of $\varphi(L'(p))$) and then solve it in time complexity $\mathcal{O}(T \log T + n(\alpha + \frac{L'(p)}{T} \log T))$ for multi-task.

To approximate optimal performance, we can just set T as \sqrt{np} and the time complexity can be revised as $\mathcal{O}(\sqrt{np} \log \sqrt{np})$ or $\mathcal{O}(\sqrt{np})$.

By the way, you can get $L(p)$ and the period of 2^k in modulo $L'(p)$ by Baby-step giant-step algorithm instead of utilizing some conclusions in number theory (if you had known or guessed that $L(p) \leq \frac{8}{3}p, :P$).

Problem H. Inner Product

If we enumerate an edge (u, v) in the first tree and remove it from the first tree, the first tree will split into two parts. And if we color the vertex in the second tree according the part which it belong to in the first tree, the contribution of (u, v) will be

$$w_1(u, v) \sum_{x=1}^n \sum_{y=1}^n [col_x \neq col_y] d_2(x, y)$$

Use the technique **dsu on tree** to maintain the color of each vertex in the second tree, and, at the same time, use **centroid decomposition** to maintain the sum of distances between every two vertices with different colors in the second tree.

The time complexity will be $\mathcal{O}(n \log^2 n)$. There also exists an $\mathcal{O}(n \log n)$ solution, we leave it to the reader.

Problem I. Counting Polygons

Formally, this problem requires you to count the number of distinct integer sequences $A = [a_0, a_1, \dots, a_{m-1}]$ meeting the following conditions:

- $\sum_{i=0}^{m-1} a_i = n$, where $a_i \in \mathbb{Z}^+$;
- $2 \cdot \max\{a_0, a_1, \dots, a_{m-1}\} < n$;
- if there exists another integer sequence $B = [b_0, b_1, \dots, b_{m-1}]$ meeting the above conditions and an integer k such that $a_i = b_{(i+k) \bmod m}$ for $i = 0, 1, \dots, m-1$, then A and B are considered the same and they should be counted only once;
- if there exists another integer sequence $B = [b_0, b_1, \dots, b_{m-1}]$ meeting the above conditions and an integer k such that $a_i = b_{m-1-i}$ for $i = 0, 1, \dots, m-1$, then A and B are considered the same;
- if A, B are considered the same, and B, C are considered the same, then A and C are considered the same.

If we ignore the second condition, we can count the number using Burnside's lemma. To deal with the second condition, we can use the Inclusion-Exclusion principle, since there may be at most one a_i that breaks the condition when $m \geq 3$. In this case, we don't need to consider all the equivalence classes.

By the way, solution with time complexity $\mathcal{O}\left(\max\{n\} + \sum_T d(\gcd(n, m))\right)$ should be accepted, where $d(x)$ is the number of divisors of x . It can get an accepted verdict because $d(x) \leq 500$ for $x \in [1, 10^7]$.

Problem J. Square Graph

First, recall the method to extract all the squares using the suffix arrays of string a and its reversed string.

Denote the half-length of the square as L , find the longest common prefix of $a[iL, n]$ and $a[iL+L, n]$, whose length is denoted by p , and the longest common suffix of $a[1, iL]$ and $a[1, iL+L]$, whose length is denoted by q . If $p+q-1 \geq L$, then we found each substring of length $2L$ obtained from $a[iL-p+1, iL+L+q-1]$ is a square.

Let's enumerate the half-length L of the square according to sorted order of w_i and find all the squares. We can get a set of triples (x, y, L) which means for $i = x, x+1, \dots, y$, i and $(i+L)$ should be connected with an edge of weight L .

Next, let's try to use two kinds of transformations which we refer to as **Split** and **Reduce** to solve this subproblem for L .

A **Split** operation transforms a single triple into an equivalent system of shorter triples. For a triple (x, y, L) , we can find an integer k such that $2^k < y - x + 1$ and $2^{k+1} \geq y - x + 1$. The triple can be split into two triples $(x, x+2^k-1, L)$ and $(y-2^k+1, y, L)$.

A **Reduce** operation removes useless triples in a set of triples with the same interval length. Since the length of all the intervals is the same, we can only care about the left endpoint x of the interval. Make a graph by connected x and $(x+L)$ and find the spanning forest of the graph and only the triples in the spanning forest are useful.

Repeat those two transformations until all the lengths of intervals equal to 1, and thus we solve this subproblem. However, L can be $1, 2, \dots, \lceil \frac{n}{2} \rceil$, so we need some modification.

For the whole problem, we can maintain a disjoint union set for every k from 0 to $\lfloor \log_2 n \rfloor$, and share the disjoint union sets to find the spanning forest. When we want to merge $(x+i)$ and $(y+i)$ ($i = 0, 1, \dots, 2^k-1$), first let's check whether x and y are already connected in dsu_k , and if so, then it is done and we can just break it. Otherwise, we need to connect x and y in dsu_k and then recursively consider $(x, y, 2^{k-1})$ and $(x+2^{k-1}, y+2^{k-1}, 2^{k-1})$ in dsu_{k-1} . It works since there are at most $\mathcal{O}(n)$ times of merging in each dsu , no matter how many L we consider.

The time complexity will be $\mathcal{O}(n \log n \alpha(n))$.

Problem K. Three Dimensions

Let $dp(i, S, T)$ be the number of pairs (a, b) and the sum of the contributions to $d(a, b)$ when

- the highest i bits of x_a, y_a, z_a, x_b, y_b and z_b are considered;
- $S = (j_1, j_2, \dots, j_9)$ is a 9-bit integer, where j_1, j_2 and j_3 mean the signs of $(x_a - x_b)$, $(y_a - y_b)$, and $(z_a - z_b)$ respectively; j_4, j_5 and j_6 mean whether $|x_a - x_b|$, $|y_a - y_b|$, and $|z_a - z_b|$ are the maximum value among them respectively; j_7, j_8 and j_9 mean whether we have to borrow a bit from the i -th highest bit to the $(i-1)$ -th highest bit in corresponding subtractions;
- $T = (k_1, k_2, \dots, k_6)$ is a 6-bit integer, where k_1, k_2, k_3, k_4, k_5 and k_6 mean whether the highest i bits of x_a, y_a, z_a, x_b, y_b and z_b are equal to the highest i bits of their own upper bounds respectively.

The recurrence equation is quite straightforward.